

Accelerating Test-Time Discovery in Complex Code Generation via Execution-Guided Self-Distillation

Mai Phong Đăng

VNUHCM University of Science
Supervisor: Ngô Minh Mẫn

2026

The problem

Test-time training (TTT): fix one hard programming problem, let the model improve itself over a few inference-time steps, then reset.

- Reward is verifiable but **sparse**: it only says *whether* an attempt passed.
- On-policy self-distillation **stalls** on hard problems — the bare student rarely produces a correct attempt to learn from.
- This is the **flat-reward trap**: no correct rollout $\Rightarrow$ no usable gradient.

*Goal: accelerate test-time **discovery** on complex code generation.*


Background: SDPO and the gap

SDPO (Hübotter et al.): extract a dense signal from the *rich feedback* that verifiable environments already produce (failing tests, runtime errors).

- The same model, conditioned on feedback, is a **self-teacher**; its per-token distribution is distilled into the student.
- Denser credit assignment than GRPO: a K -dim target per token, not one scalar per sequence.


The gap: the parent method is **student-first** — it distills the student's *own failed rollout*. On hard problems there is little that is correct to learn from.

Our idea: teacher-first



Teacher-first: distil a *correct, filtered* trajectory the teacher generates — between on-policy SDPO and off-policy SFT.

The teacher-first step




Verifier + judge filter for *correct AND independent*; the filter never distils a copy — when every candidate is a copy, the good pool is empty and that step distils nothing.

Research questions

- Method** Does teacher-first beat student-first at discovering hard code solutions — and does it survive at *matched compute*?
- RQ1** Does the *reprompt-template* formulation of the feedback affect discovery?
- RQ2** Does distillation from rich context *suppress epistemic verbalization* on code?
- RQ3** What are the *structural & domain limitations* going from procedural code to value-only math?

- Model: **Qwen3-4B** (LoRA), thinking ON; LiveCodeBench v6 (code), AIME 2026 (math pilot).
- Per problem: reset to base, **15 TTT steps**, evaluate student PRE vs POST.
- **8 frontier code problems** $\times$ **6 seeds** (medium scale) $\Rightarrow$ Wilcoxon signed-rank test.
- Metrics: pass@16, discovery@k, total generations + GPU-time.


RQ-method: teacher-first wins



Mean over 6 seeds; error bars = 1 SD. Wilcoxon signed-rank $p < 0.01$.

- TF > SF on every problem (mean), largest margins on the hardest — **39 / 8 / 1** over 48 seed-pairs.
- Wilcoxon $p < 0.01$ (8 problems $\times$ 6 seeds) — not a crude sign test.
- **RQ1**: on hard problems, template order **T5 > T2 > T1** (reasoning-inducing best).


The headline: escape-zero



Student-first stays flat at 0; teacher-first escapes from step ~4.

- Student-first stays **flat at zero**.
- Teacher-first **lifts off** from $\sim$ step 4.
- Direct signature of the flat-reward trap: on-policy has nothing correct; off-policy injection does.


Compute-fair: not just more sampling (RQ-method)



Compute-to-correct frontier: TF reaches a target discovery at ~2x fewer generations.

- Match the budget: student-first at **10 generations/step**.
- Teacher-first still wins, and reaches a target discovery with **$\sim 2\times$ fewer total generations**.
- The advantage is the *kind* of distilled trajectory, not sample volume.


RQ2: epistemic suppression on code



Epistemic markers decline and accumulate downward across TTT steps (code, thinking ON).

- Uncertainty markers (“wait”, “let me reconsider”) **decline** across steps, accumulating.
- Connects to Kim et al.: high $I(y; c | x) \Rightarrow$ suppression.
- *On code* this coincides with **improved** correctness — consolidation, not answer-copying.

RQ3: value vs. procedure — when it works



- Code output *is* a method (a program that generalizes) $\Rightarrow$ distillation installs a procedure $\Rightarrow$ **escape**.
- Answer-only math output is a value $\Rightarrow$ the teacher can only copy $\Rightarrow$ **no escape**.
- Escape needs an **in-reach method**, not merely a correct value.

- ① A **teacher-first** organization of execution-guided self-distillation (clean filter, never distils a copy).
- ② Teacher-first **significantly accelerates discovery** and **escapes the flat-reward trap** — compute-fair (RQ-method).
- ③ A **reprompt-template effect**, strongest on hard problems (RQ1).
- ④ A measured **epistemic-suppression** result on code (RQ2).
- ⑤ A **value-versus-procedure boundary** characterizing *when* the method helps.

Limitations & publication roadmap

Honest bounds (§6.2):

- One **4B** Qwen family at a personal compute scale — a characterization, not a scaling law.
- Math pilot bottlenecked by **AIME**'s answer-only feedback (no worked solutions to derive from).
- A stronger base (**8B+**) is expected to *narrow* the TF-vs-SF margin on reachable tasks.
- Epistemic result tied to one model family / lexical marker set.

Roadmap (post-graduation): method-bearing math references (**MATH-500**) to lift the math bottleneck; scale to **8B+** and a second code benchmark; broaden the epistemic measurement across tokenizers and scales.

- **Teacher-first execution-guided self-distillation** accelerates test-time discovery in complex code generation.
- It escapes the flat-reward trap, the advantage is **compute-fair**, and it carries a measured epistemic effect.
- The **value-versus-procedure boundary** explains *when* it works: an in-reach method, not just an answer.

Thank you.

Questions?